

---

# CIS 471/571 (Winter 2020): Introduction to Artificial Intelligence

## Lecture 6: Adversarial Search

---

Thanh H. Nguyen

Source: <http://ai.berkeley.edu/home.html>



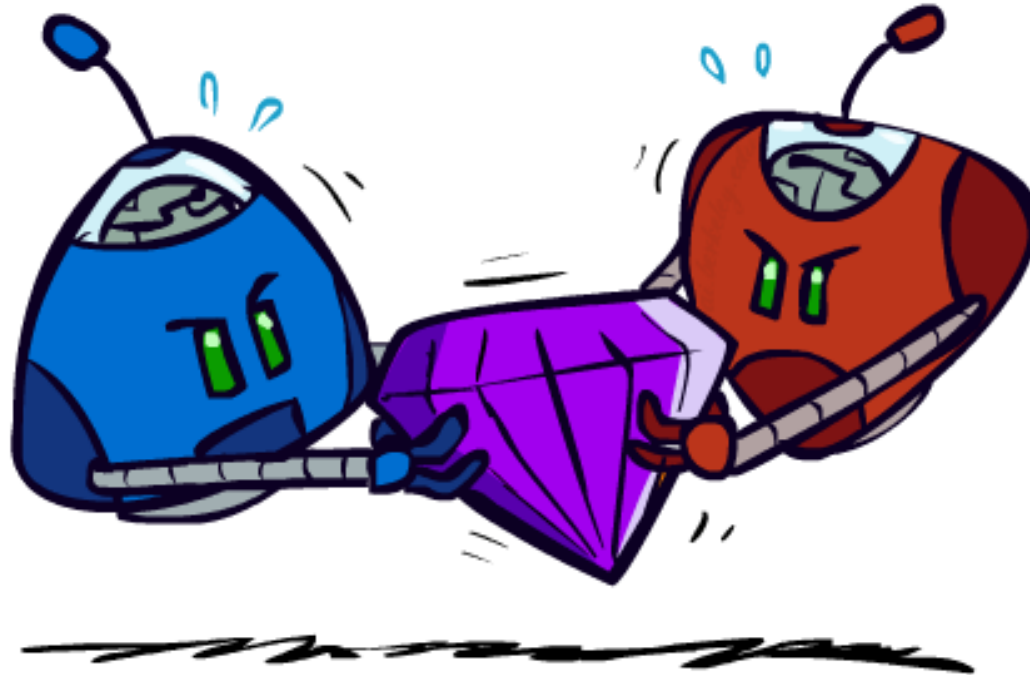
# Announcements

---

- Project 2:
  - Deadline: Feb 02, 2020
- Homework 2:
  - Deadline: Feb 03, 2020


# Adversarial Games

---




# Types of Games

- Many different kinds of games!
- Axes:
  - Deterministic or stochastic?
  - One, two, or more players?
  - Zero sum?
  - Perfect information (can you see the state)?
- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state




# Deterministic Games

- Many possible formalizations, one is:
  - States:  $S$  (start at  $s_0$ )
  - Players:  $P=\{1...N\}$  (usually take turns)
  - Actions:  $A$  (may depend on player / state)
  - Transition Function:  $S \times A \rightarrow S$
  - Terminal Test:  $S \rightarrow \{t, f\}$
  - Terminal Utilities:  $S \times P \rightarrow R$
- Solution for a player is a **policy**:  $S \rightarrow A$



# Zero-Sum Games



- Zero-Sum Games


- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition

- General Games


- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible
- More later on non-zero-sum games

# Adversarial Search

---




# Single-Agent Trees





# Value of a State

Value of a state:  
The best achievable  
outcome (utility)  
from that state



Non-Terminal States:


$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

Terminal States:

$$V(s) = \text{known}$$



# Adversarial Game Trees




# Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$




Terminal States:

$$V(s) = \text{known}$$




# Tic-Tac-Toe Game Tree



# Adversarial Search (Minimax)

- Deterministic, zero-sum games:
  - Tic-tac-toe, chess, checkers
  - One player maximizes result
  - The other minimizes result
- Minimax search:
  - A state-space search tree
  - Players alternate turns
  - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



# Minimax Implementation

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)


def max-value(state):

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return  $v$



def min-value(state):

initialize  $v = +\infty$

for each successor of state:

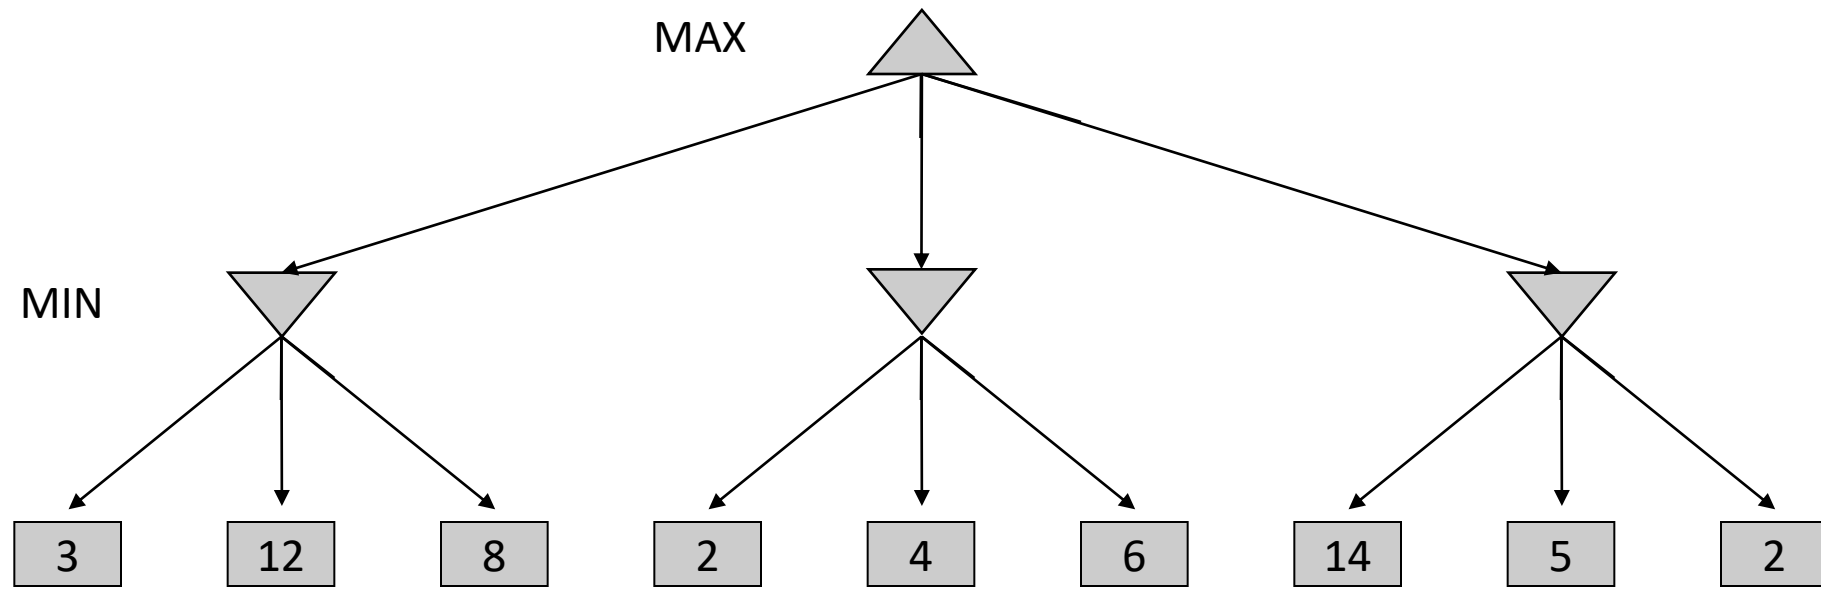
$v = \min(v, \text{value}(\text{successor}))$

return  $v$






# Minimax Example

---



# Minimax Properties




Optimal against a perfect player. Otherwise?





# Minimax Efficiency

- How efficient is minimax?
  - Just like (exhaustive) DFS
  - Time:  $O(b^m)$
  - Space:  $O(bm)$
- Example: For chess,  $b \approx 35$ ,  $m \approx 100$ 
  - Exact solution is completely infeasible
  - But, do we need to explore the whole tree?




# Resource Limits

---



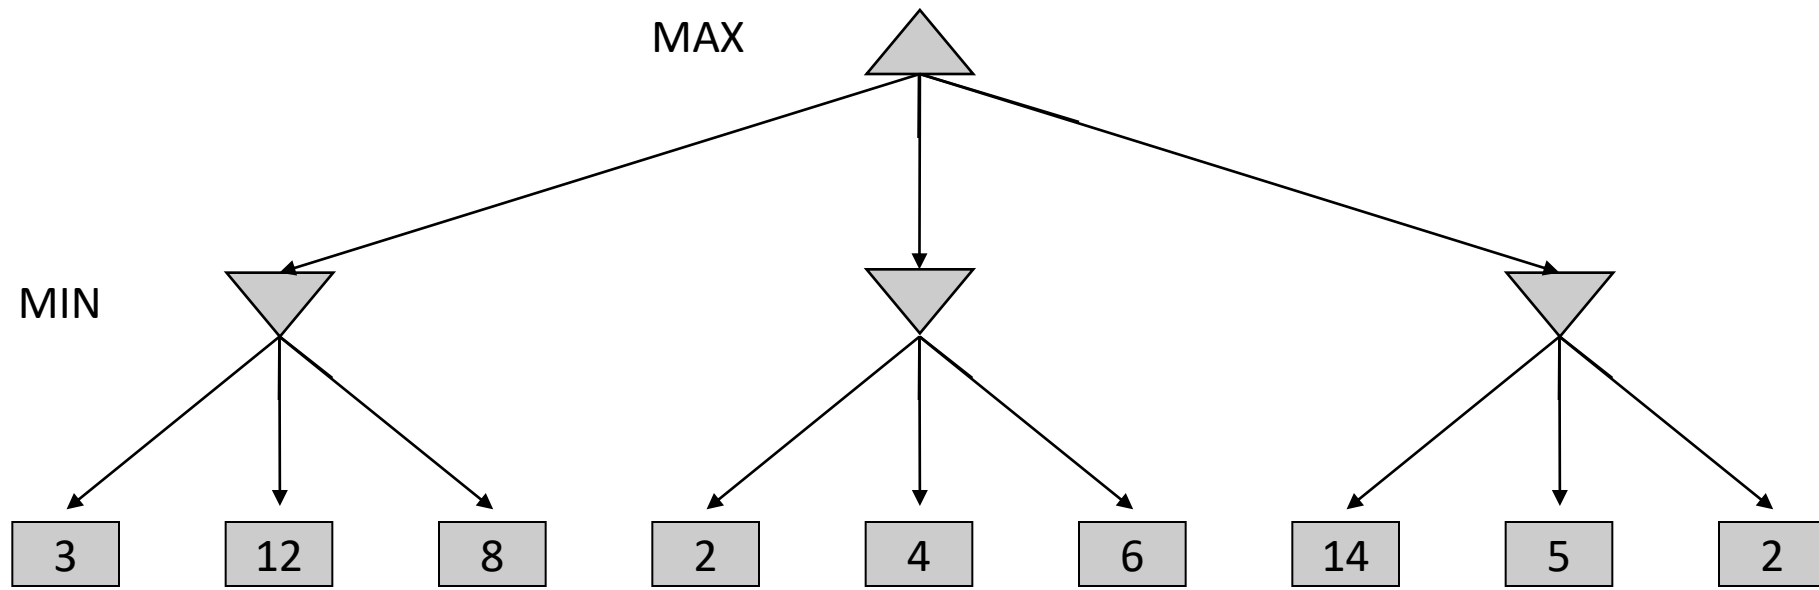
# Game Tree Pruning

---




# Minimax Example

---



# Minimax Pruning

---



# Alpha-Beta Pruning

- **Alpha  $\alpha$** : value of the best choice so far for MAX (lower bound of Max utility)
- **Beta  $\beta$** : value of the best choice so far for MIN (upper bound of Min utility)
- Expanding at MAX node **n**: update  $\alpha$ 
  - If a child of **n** has value greater than  $\beta$ , stop expanding the MAX node **n**
  - Reason: MIN parent of **n** would not choose the action which leads to **n**
- At MIN node **n**: update  $\beta$ 
  - If a child of **n** has value less than  $\alpha$ , stop expanding the MIN node **n**
  - Reason: MAX parent of **n** would not choose the action which leads to **n**

# Alpha-Beta Implementation

def value(state,  $\alpha$ ,  $\beta$ ):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state,  $\alpha$ ,  $\beta$ )

if the next agent is MIN: return min-value(state,  $\alpha$ ,  $\beta$ )

def max-value(state,  $\alpha$ ,  $\beta$ ):

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if  $v \geq \beta$  return  $v$

$\alpha = \max(\alpha, v)$

return  $v$

def min-value(state,  $\alpha$ ,  $\beta$ ):

initialize  $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


if  $v \leq \alpha$  return  $v$

$\beta = \min(\beta, v)$

return  $v$

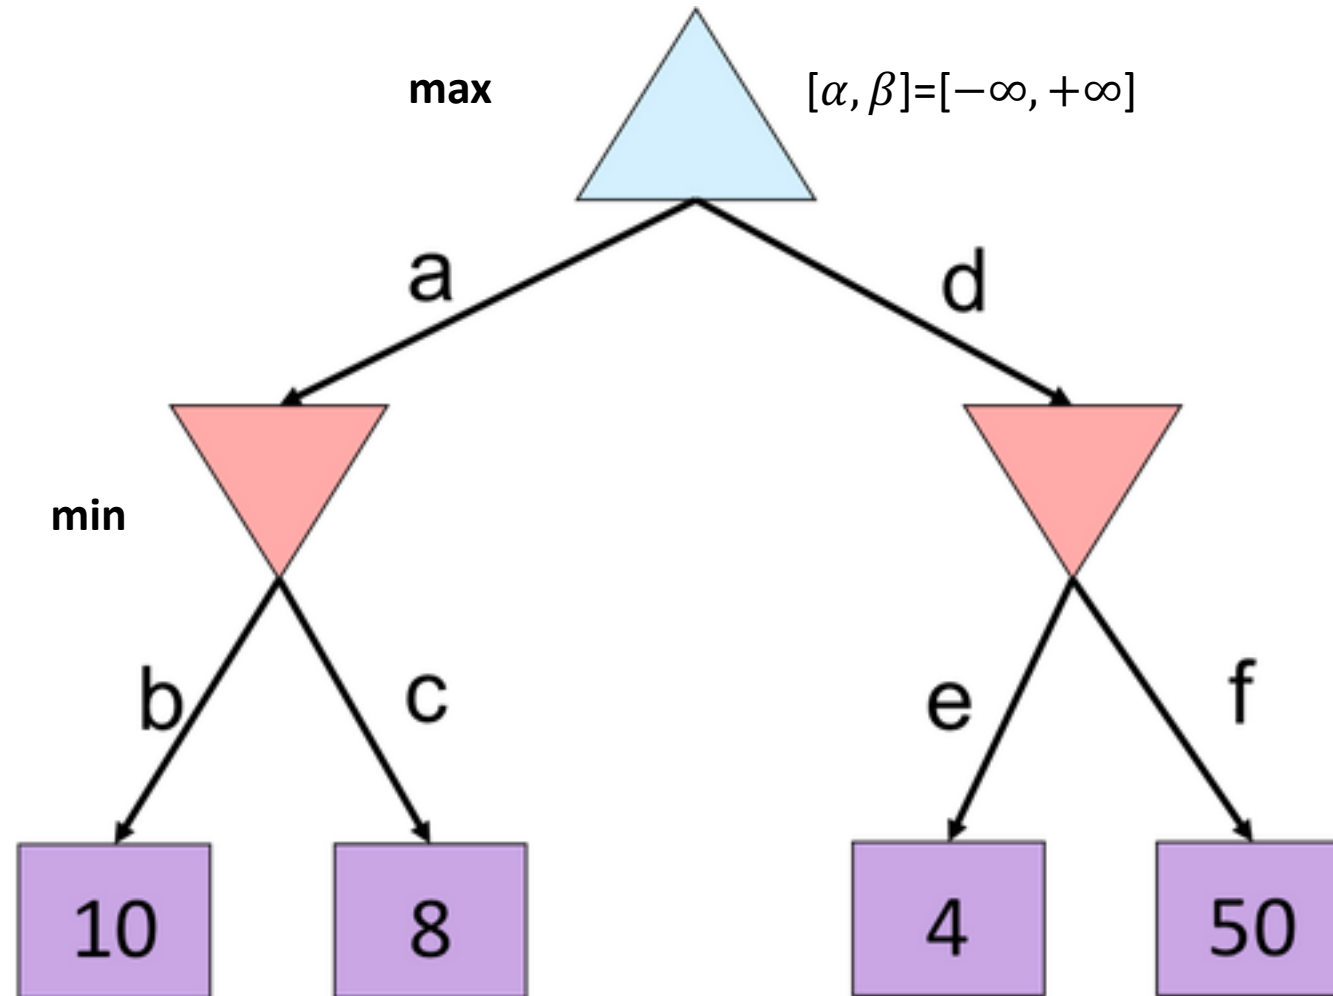
# Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong
  - Important: children of the root may have the wrong value
  - So the most naïve version won't let you do action selection
- Good child ordering improves effectiveness of pruning







# Alpha-Beta Quiz




# Alpha-Beta Quiz




# Alpha-Beta Quiz



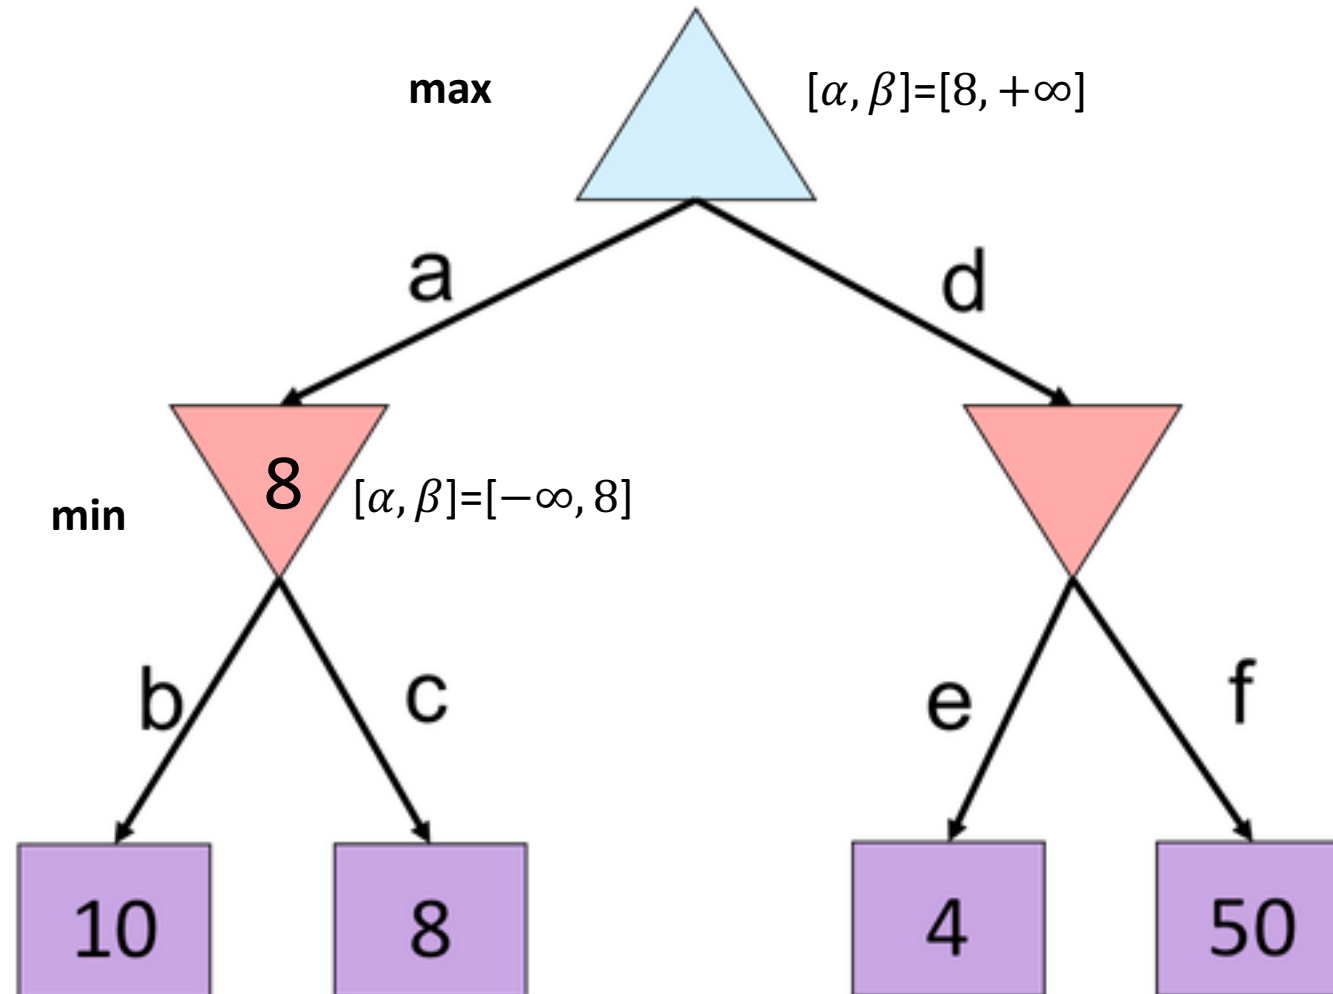
# Alpha-Beta Quiz



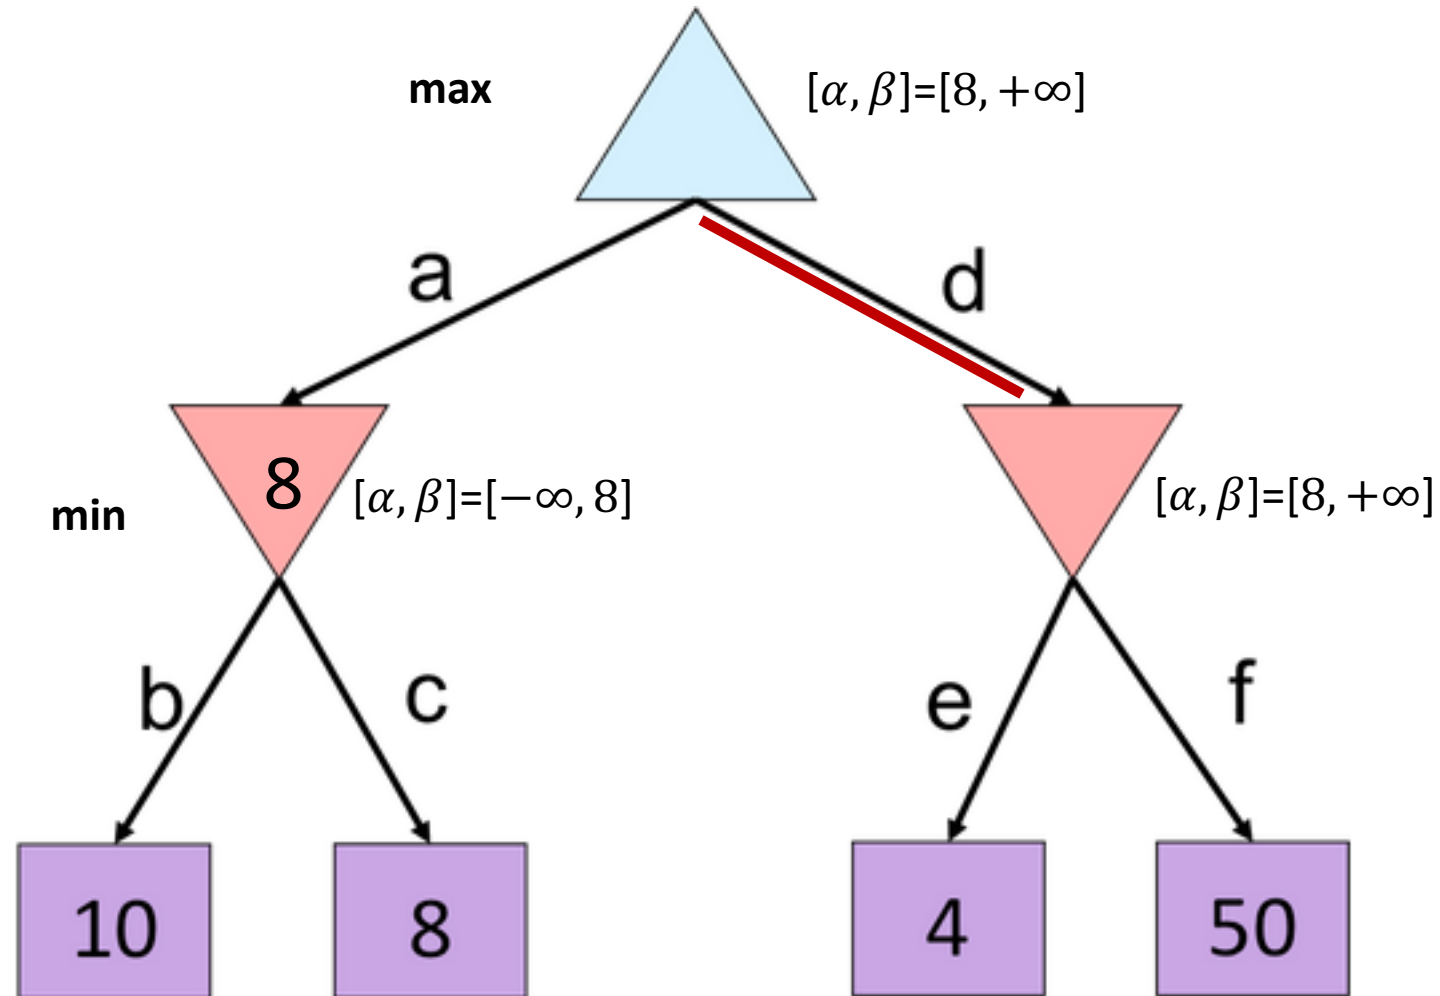
# Alpha-Beta Quiz




# Alpha-Beta Quiz



# Alpha-Beta Quiz




# Alpha-Beta Quiz

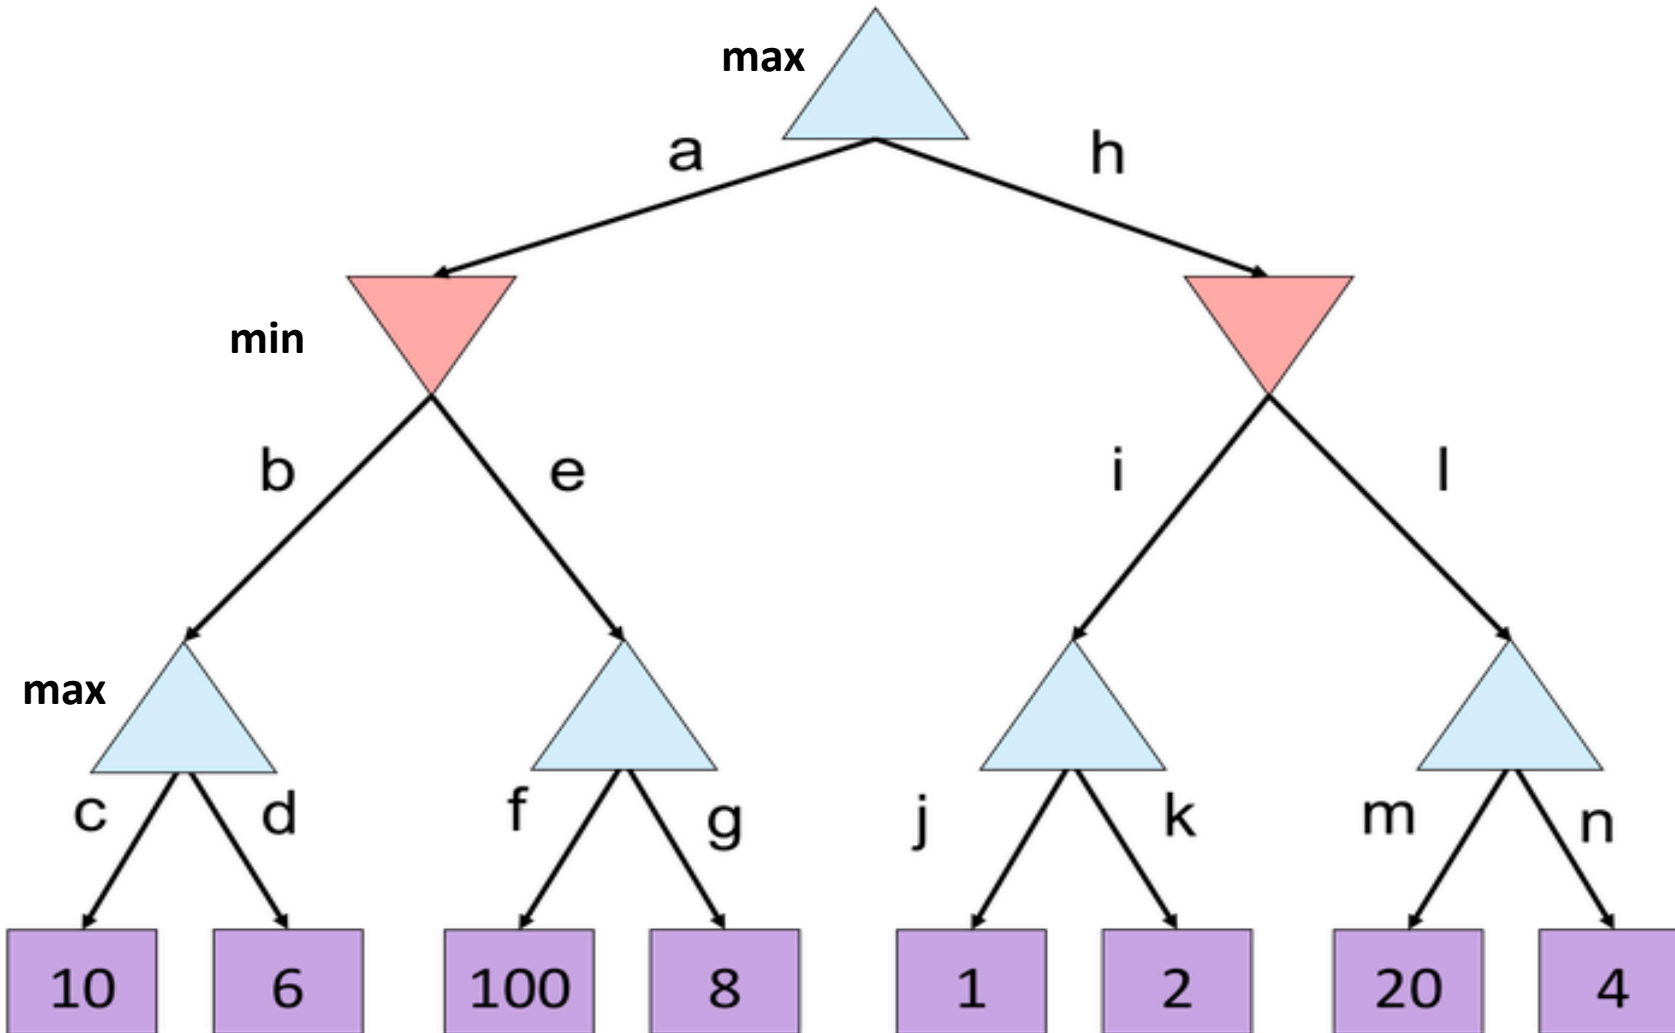





# Alpha-Beta Quiz




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2



# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2







# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2



# Alpha-Beta Quiz 2




# Alpha-Beta Quiz 2

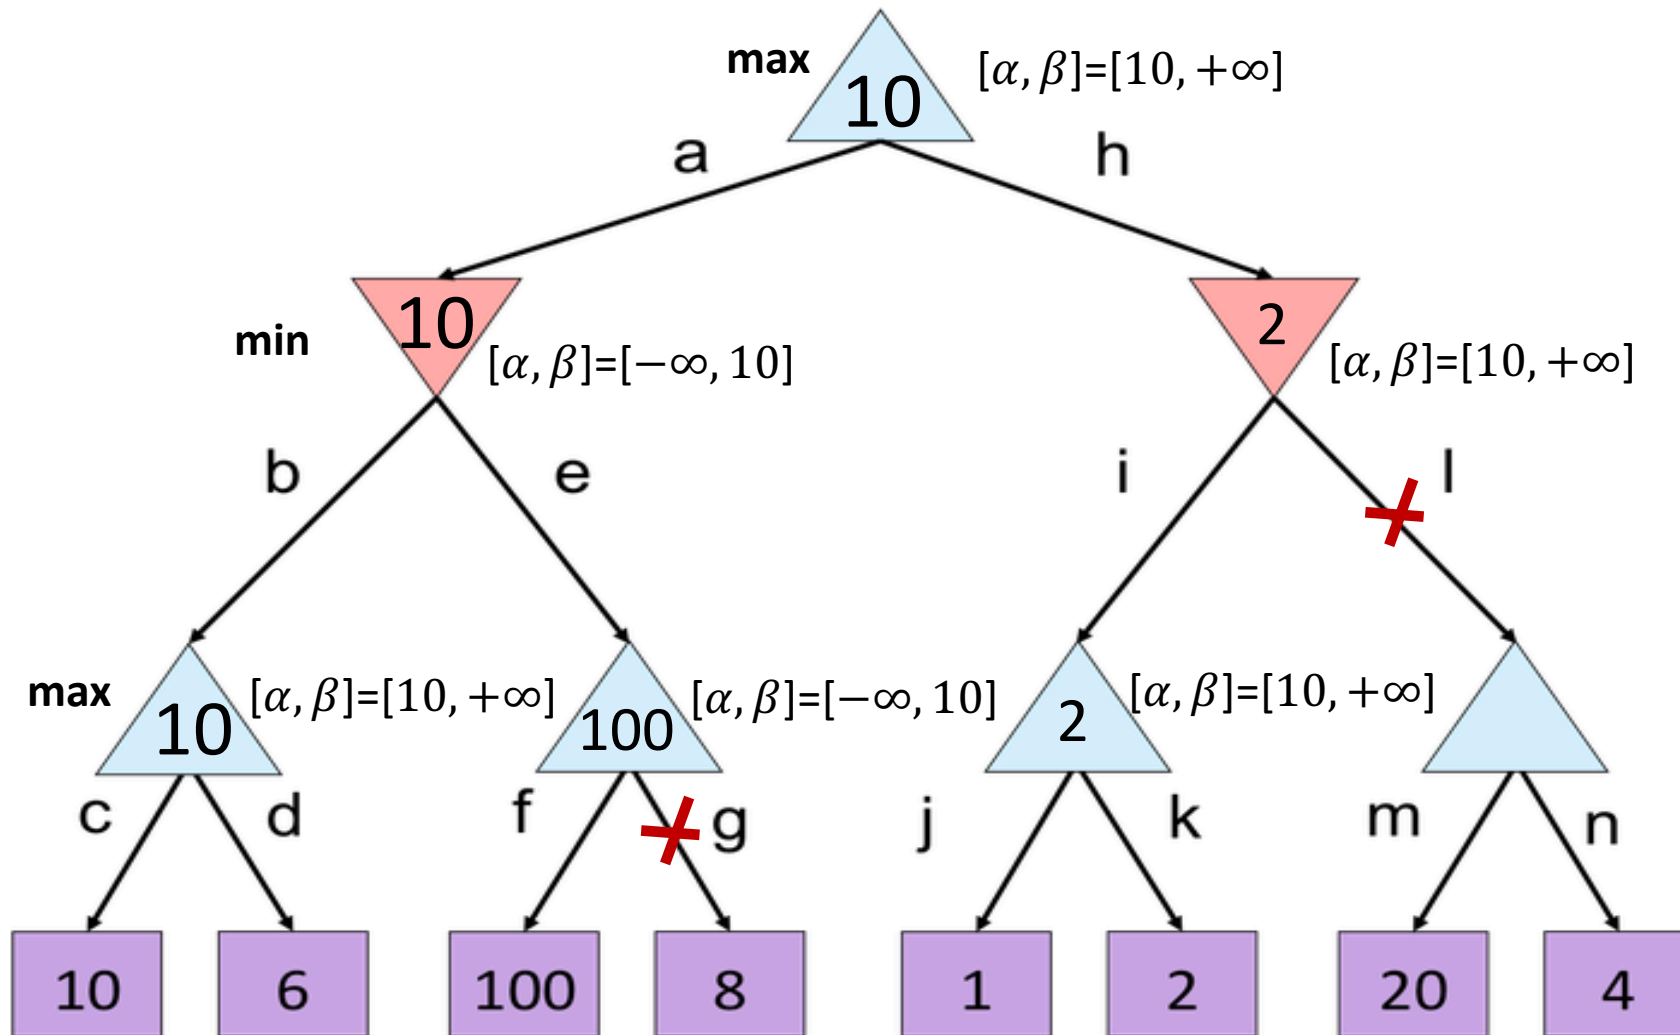




# Alpha-Beta Quiz 2



# Alpha-Beta Quiz 2




# Resource Limits

---




# Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
  - Instead, search only to a limited depth in the tree
  - Replace terminal utilities with an evaluation function for non-terminal positions
- Example:
  - Suppose we have 100 seconds, can explore 10K nodes / sec
  - So can check 1M nodes per move
  - $\alpha$ - $\beta$  reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm



# Why Pacman Starves




- A danger of replanning agents!
  - He knows his score will go up by eating the dot now (west, east)
  - He knows his score will go up just as much by eating the dot later (east, west)
  - There are no point-scoring opportunities after eating the dot (within the horizon, two here)
  - Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!




# Evaluation Functions

---



# Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ , etc.



# Evaluation for Pacman



---





# Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation



# Synergies between Evaluation Function and Alpha-Beta?

---

- Alpha-Beta: amount of pruning depends on expansion ordering
  - Evaluation function can provide guidance to expand most promising nodes first (which later makes it more likely there is already a good alternative on the path to the root)
    - (somewhat similar to role of A\* heuristic, CSPs filtering)
- Alpha-Beta: (similar for roles of min-max swapped)
  - Value at a min-node will only keep going down
  - Once value of min-node lower than better option for max along path to root, can prune
  - Hence: IF evaluation function provides upper-bound on value at min-node, and upper-bound already lower than better option for max along path to root THEN can prune

